The purpose of this note is to supplement the material on linear
regression in the textbook.

Regression
==========

Given two lists of numbers X,Y of the same length and a linear function
y=L(x)=ax+b (a and b are constants), the **rms error for the line L** is the
rms of the list Y-L(X). (For example, if A is an entry in the list X and
B is the corresponding entry in the list Y, then the corresponding entry
in the list Y-L(X) is B-(aA+b).) We interpret the rms error for a line
as a measure of how well L performs as a tool for predicting the Y value
that corresponds to a given X value. Given an X value A, the value
B=L(A) has a "prediction error" equal to B-L(A). The rms of the list of
prediction errors is low if L is a good prediction tool and high if L is
a poor prediction tool. We interpret the rms error for a line L as the
size of the typical error that we would observe if we used L to predict
an unknown Y value for a given X value.

Among all possible lines, there is one single line with the least
possible rms error, called the **regression line** for Y on X. This means
that the regression line is the best possible line to use as a
prediction tool to guess an unknown corresponding Y value for a given X
value. The regression line passes through the point (AVE(X),AVE(Y)),
called the **point of averages** for the lists X,Y, and has slope
r*SD(Y)/SD(X), where r is the number

  r = AVE([X-AVE(X)/SD(X)]*[Y-AVE(Y)/SD(Y)])

called the **correlation coefficient** for the pair of lists X,Y. The rms
error for the regression line is

  rms error for regression line = SD(Y) * SQRT(1-r^2).

The list (X,Y) of pairs of points from the lists X,Y is called the
scatter diagram of the lists X,Y. The scatter diagram is said to show
linear association if the points are scattered roughly symmetrically
about a linear axis. The size of the correlation coefficient determines
the tightness or looseness of scattering of points (X,Y) about the
regression line. For |r| near 1, points in the scatter diagram are
tightly clustered about the regression line. As |r| approaches 0, the
scatter diagram is more and more loosely clustered about the regression
line.

A **thin vertical strip** in the scatter diagram is a set of points (X,Y)
whose X-coordinate lies in a range from A-e to A+e, where A is a number
within the range of entries of the list X and e is a small increment. If
the various sets of Y values in thin vertical strips (for various values
of A) has a similar SD as A varies, the scatter diagram is said to be
homoscedastic. Otherwise, the scatter diagram is said to be
heteroscedastic.

Let's write REGR(A) for the y-coordinate of a point on the regression
line for the x-value A. Here's a formula.

  REGR(A) = r * [A-AVE(X)/SD(X)] * SD(Y) + AVE(Y)

A **residual** is another name for a vertical error Y-REGR(X) for a data
point (X,Y). The **graph of residuals** is the set of points
(X,Y-REGR(X)). You can see linear association and homoscedasticity (or
the lack of either one) very plainly in the graph of residuals: data
lists X,Y have a linear association if the graph of residuals is roughly
symmetric across the X-axis; and the X,Y data are homoscedastic if the
graph of residuals has fairly even densities and vertical ranges in thin
vertical strips across the horizontal range of X.

For the sake of comparison, here are three lines and their rms
errors. All three lines pass through the point of averages. The SD line
is defined to have positive slope if r>0 and negative slope if r<0.

line	slope	rms error
===============	==============	===================
regression line	r*SD(Y)/SD(X)	SD(Y) * SQRT(1-r^2)
SD line	+/- SD(Y)/SD(X)	SD(Y) * SQRT(2(1-|r|))
average line	0	SD(Y)

```
Facts about the correlation coefficient r
=========================================
* The value of r is in the range -1 to 1. The value of |r| is 1
  precisely when all the points (X,Y) on the scatter diagram lie on a
  straight line.

* The correlation coefficient is unitless. The value of r does not
  change if you rescale X or Y.

* It is tempting to summarize data using averages, but this has a
  misleading effect on correlation. Suppose that the points in a scatter
  diagram are subdivided into subsets, and that each subset is replaced
  by a single point

    (average of the X values in the subset, average of the Y values in the subset).

  The correlation for this smaller data set of average points is usually
  higher (in absolute value) than the correlation for the original
  scatter diagram. This phenomenon is called **_ecological correlation_**.

* A strong correlation provides no evidence of a causal link between the
  variables X and Y.

The regression effect and regression fallacy
============================================
Start with an X-value A, somewhere above or below the value AVE(X). Use
the regression line for Y on X to predict the average Y value for data
points whose X value is near A. Call this prediction B. Now use the
regression line for X on Y to predict the average X value for data
points whose Y value is near B. Call this prediction C. It will always
turn out that C is between AVE(X) and A. This is due simply to the fact
that, for real data, the value of |r| is less than 1. This is called the
**_regression effect_**. Any attempt to explain why C is between AVE(X) and A
by any other reason than simply the fact that |r|<1 is called a
**_regression fallacy_**.
```